

# CS251/CS253: COMPUTING TOOLS STRUCTURED QUERY LANGUAGE (SQL)

**Arnab Bhattacharya**

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

5th February, 2020

# Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases

# Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
  - Can access and manipulate data stored as a particular data model

# Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
  - Can access and manipulate data stored as a particular data model
- *Declarative language*
  - Specifies what to do, but not how to do

# Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
  - Can access and manipulate data stored as a particular data model
- *Declarative language*
  - Specifies what to do, but not how to do
- Is also a **data definition language (DDL)**
  - Defines database relations and schemas

# Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
  - Can access and manipulate data stored as a particular data model
- *Declarative language*
  - Specifies what to do, but not how to do
- Is also a **data definition language (DDL)**
  - Defines database relations and schemas
- SQL has evolved widely after its first inception
  - Supports lots of extra operations that are non-standard

# Relational Algebra

- SQL is based upon relational algebra
- Relational data model
- Operators are functions from one or two input relations to an output relation

# Relational Algebra

- SQL is based upon relational algebra
- Relational data model
- Operators are functions from one or two input relations to an output relation
- Uses *propositional calculus* formed by *expressions* connected by
  - 1 and:  $\wedge$
  - 2 or:  $\vee$
  - 3 not:  $\neg$
- Each term is of the form  
`<attr/const> comparator <attr/const>`  
where comparator is one of  $=, \neq, >, \geq, <, \leq$



# Relation

- A **relation** is a subset of the cross-product of sets
- For sets  $D_1, D_2, \dots, D_n$ , a relation  $r$  is a set of  $n$ -ary tuples of the form  $(a_1, a_2, \dots, a_n)$  where each  $a_i \in D_i$
- Example
  - `name = {A, B, C}`
  - `designation = {L, E, W}`
  - `identifier =  $\mathcal{N}$`
  - $r = \{(A, E, 4), (B, E, 9), (C, W, 23)\}$  is a relation over  
`name  $\times$  designation  $\times$  identifier`
- Relations are *unordered*

# Relation

- A **relation** is a subset of the cross-product of sets
- For sets  $D_1, D_2, \dots, D_n$ , a relation  $r$  is a set of  $n$ -ary tuples of the form  $(a_1, a_2, \dots, a_n)$  where each  $a_i \in D_i$
- Example
  - `name` = {A, B, C}
  - `designation` = {L, E, W}
  - `identifier` =  $\mathcal{N}$
  - $r = \{(A, E, 4), (B, E, 9), (C, W, 23)\}$  is a relation over  
`name`  $\times$  `designation`  $\times$  `identifier`
- Relations are *unordered*
- Generally depicted as a table

name	designation	identifier
A	E	4
B	E	9
C	W	23

# Attribute

- Each attribute of a relation has a name
- There is a domain for each attribute
- Attributes are *generally* **atomic**
  - Indivisible, not sets
- Domain is atomic if all members are atomic
- Special value **null** in every domain

# Relation Schema and Tuple

- The sets define a **relation schema**
- Example
  - Schema is `Person = (name, designation, identifier)`
- Relations are defined over a schema
- If schema is  $R$ , relation is denoted by  $r(R)$ 
  - Example: `persons(Person)`
- A **relation instance** is a particular instance from the schema
  - Earlier example
- An element of a relation (instance) is a **tuple**

# Relation Schema and Tuple

- The sets define a **relation schema**
- Example
  - Schema is `Person = (name, designation, identifier)`
- Relations are defined over a schema
- If schema is  $R$ , relation is denoted by  $r(R)$ 
  - Example: `persons(Person)`
- A **relation instance** is a particular instance from the schema
  - Earlier example
- An element of a relation (instance) is a **tuple**
- Tuples are rows and attributes are columns

# Database

- Consists of multiple *inter-related* relations
- Each relation stores information about a particular relationship

# Database

- Consists of multiple *inter-related* relations
- Each relation stores information about a particular relationship
- Alternatively, a single relation can store all data
- Problems
  - Data repetition
  - Need for null values
- Normalization theory deals with how to design relation schemas

# Key

- $K \subseteq R$  is a **superkey** of  $R$  if and only if values for  $K$  are sufficient to identify a unique tuple in *all* possible relations  $r(R)$ 
  - Possible  $r(R)$  signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
  - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey



# Key

- $K \subseteq R$  is a **superkey** of  $R$  if and only if values for  $K$  are sufficient to identify a unique tuple in *all* possible relations  $r(R)$ 
  - Possible  $r(R)$  signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
  - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey
- A superkey  $K$  is a **candidate key** if  $K$  is minimal, i.e., no proper subset of it is a superkey
  - {name} is a candidate key

# Key

- $K \subseteq R$  is a **superkey** of  $R$  if and only if values for  $K$  are sufficient to identify a unique tuple in *all* possible relations  $r(R)$ 
  - Possible  $r(R)$  signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
  - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey
- A superkey  $K$  is a **candidate key** if  $K$  is minimal, i.e., no proper subset of it is a superkey
  - {name} is a candidate key
- There may be multiple candidate keys
  - {name}, {identifier} are candidate keys

# Key

- $K \subseteq R$  is a **superkey** of  $R$  if and only if values for  $K$  are sufficient to identify a unique tuple in *all* possible relations  $r(R)$ 
  - Possible  $r(R)$  signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
  - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey
- A superkey  $K$  is a **candidate key** if  $K$  is minimal, i.e., no proper subset of it is a superkey
  - {name} is a candidate key
- There may be multiple candidate keys
  - {name}, {identifier} are candidate keys
- **Primary key** is a candidate key chosen to serve as the primary means of identifying tuples
  - Choice is arbitrary as it depends on the database designer
  - Other candidate keys are called **secondary keys**

# Foreign Key

- A relation schema may have an attribute that is unique (e.g., a primary key) in another schema
- This attribute is then called a **foreign key**
- Example
  - depositor = (name, number)
  - customer = (name, street, city)
  - account = (number, balance)
  - name and number in depositor are foreign keys
- Values in the foreign key attribute of the **referencing relation** may only come from those in the primary key of the **referenced relation**

# Example Schema

- course (code, title, *ctype*, webpage)
- coursetype (ctype, *dept*)
- faculty (fid, name, *dept*, designation)
- department (deptid, name)
- semester (yr, half)
- offering (*coursecode*, yr, half, instructor)
- student (roll, name, *dept*, cpi)
- program (roll, *p*type)
- registration (*coursecode, roll, yr, half, grade*code)
- grade (gradecode, value)

# Creating Relation Schemas

- **create table**: create table

$r(A_1 D_1 C_1, \dots, A_n D_n C_n, (IC_1), \dots, (IC_k))$

- $r$  is the name of the relation
- Each  $A_i$  is an attribute name whose data type or domain is specified by  $D_i$
- $C_i$  specifies constraints or settings (if any)
- $IC_j$  represents integrity constraints (if any)

- Example

```
create table faculty (  
    fid integer primary key,  
    name varchar(50) not null,  
    dept integer,  
    designation varchar(3)  
);
```

# Data Types in SQL

- *char(n)*: fixed-length character string
- *varchar(n)*: variable-length character string, up to  $n$
- *integer* or *int*: integer
- *smallint*: short integer
- *numeric(n,d)*: floating-point number with a total of  $n$  digits of which  $d$  is after the decimal point
- *real*: single-precision floating-point number
- *double precision*: double-precision floating-point number
- *float(n)*: floating-point number with at least  $n$  digits

# Data Types in SQL

- *char(n)*: fixed-length character string
- *varchar(n)*: variable-length character string, up to  $n$
- *integer* or *int*: integer
- *smallint*: short integer
- *numeric(n,d)*: floating-point number with a total of  $n$  digits of which  $d$  is after the decimal point
- *real*: single-precision floating-point number
- *double precision*: double-precision floating-point number
- *float(n)*: floating-point number with at least  $n$  digits
- *date*: yyyy-mm-dd format
- *time*: hh:mm:ss format
- *time(i)*: hh:mm:ss.i . . . i format with additional  $i$  digits for fraction of a second
- *timestamp*: both date and time
- *interval*: relative value in either year-month or day-time format



# Other Data Types

- Large objects such as images, videos, strings can be stored as
  - **blob**: binary large object
  - **clob**: character large object
  - A pointer to the object is stored in the relation, and not the object itself

# Constraints

- Can be specified for each attribute as well as separately
  - *not null*: the attribute cannot be null
    - Requires some value while inserting as otherwise null is the default
  - *primary key* ( $A_i, \dots, A_j$ ): automatically ensures not null
  - *default n*: defaults to  $n$  if no value is specified
  - *unique*: specifies that this is a candidate key
  - *foreign key*: specifies as a foreign key and the relation it refers to
  - *check P*: predicate  $P$  must be satisfied

```
create table faculty (  
    fid integer ,  
    name varchar(50) not null ,  
    dept integer ,  
    designation varchar(3) ,  
    primary key fid ,  
    foreign key (dept) references department ,  
    check (fid >= 0)  
);
```

# Deleting or Modifying a Relation Schema

- **drop table**: **drop table** *r* deletes the table from the database
  - Must satisfy other constraints already applied
- Example
  - drop table** faculty;

# Deleting or Modifying a Relation Schema

- **drop table**: **drop table**  $r$  deletes the table from the database
  - Must satisfy other constraints already applied

- Example

```
drop table faculty;
```

- **alter table**: **alter table**  $r$  **add**  $A$   $D$   $C$

- Adds attribute  $A$  with data type  $D$  at the end
- $C$  specifies constraints on  $A$  (if any)
- Must satisfy other constraints already applied

- **alter table**: **alter table**  $r$  **drop**  $A$

- Deletes attribute  $A$  from all tuples
- Must satisfy other constraints already applied

- Example

```
alter table faculty add room varchar(10);  
alter table course drop webpage;
```

# Basic Query Structure

- SQL is based on *relational algebra*
- A *basic* SQL query is of the form  
    select  $A_1, \dots, A_n$   
    from  $r_1, \dots, r_m$   
    where  $P$
- Each  $r_i$  is a relation
- Each  $A_j$  is an attribute from one of  $r_1, \dots, r_m$
- $P$  is a predicate involving attributes and constants
- **where** can be left out, which then means *true*
- Result is a relation with the schema  $(A_1, \dots, A_n)$

# Basic Query Structure

- SQL is based on *relational algebra*
- A *basic* SQL query is of the form  
    select  $A_1, \dots, A_n$   
    from  $r_1, \dots, r_m$   
    where  $P$
- Each  $r_i$  is a relation
- Each  $A_j$  is an attribute from one of  $r_1, \dots, r_m$
- $P$  is a predicate involving attributes and constants
- *where* can be left out, which then means *true*
- Result is a relation with the schema  $(A_1, \dots, A_n)$
- Is equivalent to the relational algebra query

# Basic Query Structure

- SQL is based on *relational algebra*
- A *basic* SQL query is of the form  
    select  $A_1, \dots, A_n$   
    from  $r_1, \dots, r_m$   
    where  $P$
- Each  $r_i$  is a relation
- Each  $A_j$  is an attribute from one of  $r_1, \dots, r_m$
- $P$  is a predicate involving attributes and constants
- *where* can be left out, which then means *true*
- Result is a relation with the schema  $(A_1, \dots, A_n)$
- Is equivalent to the relational algebra query  
     $\Pi_{A_1, \dots, A_n}(\sigma_P(r_1 \times \dots \times r_m))$

# Multisets

- SQL relations are **multisets** or **bags** of tuples and not sets
- Consequently, there may be two identical tuples
- This is the biggest distinction with relational algebra



# Multisets

- SQL relations are **multisets** or **bags** of tuples and not sets
- Consequently, there may be two identical tuples
- This is the biggest distinction with relational algebra
- The set behavior can be enforced by the keyword **unique**
- In a query, keyword **distinct** achieves the same effect
- Opposite is keyword **all**, which is *default*

# Project in Relational Algebra

- $\Pi_{A_1, \dots, A_k}(r)$
- $A_i$ , etc. are attributes of  $r$
- Select only the specified attributes  $A_1, \dots, A_k$  from all tuples of  $r$
- Duplicate rows are removed, since relations are sets

# Project in Relational Algebra

- $\Pi_{A_1, \dots, A_k}(r)$
- $A_i$ , etc. are attributes of  $r$
- Select only the specified attributes  $A_1, \dots, A_k$  from all tuples of  $r$
- Duplicate rows are removed, since relations are sets
- Schema is

# Project in Relational Algebra

- $\Pi_{A_1, \dots, A_k}(r)$
- $A_i$ , etc. are attributes of  $r$
- Select only the specified attributes  $A_1, \dots, A_k$  from all tuples of  $r$
- Duplicate rows are removed, since relations are sets
- Schema is changed
- Applying  $\Pi_{A,C}$  on

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8

returns

# Project in Relational Algebra

- $\Pi_{A_1, \dots, A_k}(r)$
- $A_i$ , etc. are attributes of  $r$
- Select only the specified attributes  $A_1, \dots, A_k$  from all tuples of  $r$
- Duplicate rows are removed, since relations are sets
- Schema is changed
- Applying  $\Pi_{A,C}$  on

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8

	A	C
returns	1	5
	2	5
	2	8

# Select

- Lists attributes in the final output

# Select

- Lists attributes in the final output
- Example: Find codes of courses offered in 2018

# Select

- Lists attributes in the final output
- Example: Find codes of courses offered in 2018

```
select coursecode  
from offering  
where yr = 2018;
```

- Case-insensitive
- **select** \* chooses all attributes
- To eliminate duplicates, use **select distinct** ...
- Otherwise, by default is **select all** ...



# Select

- Lists attributes in the final output
- Example: Find codes of courses offered in 2018

```
select coursecode  
from offering  
where yr = 2018;
```

- Case-insensitive
- **select** \* chooses all attributes
- To eliminate duplicates, use **select distinct ...**
- Otherwise, by default is **select all ...**
- Can contain arithmetic expressions

```
select coursecode, yr - 1959  
from offering  
where yr = 2018;
```

# Cartesian Product in Relational Algebra

- $r \times s = \{\langle u, v \rangle | u \in r \text{ and } v \in s\}$
- Attributes of relations  $r$  and  $s$  should be disjoint
- If attributes are not disjoint, renaming should be used
- Schema is

# Cartesian Product in Relational Algebra

- $r \times s = \{\langle u, v \rangle | u \in r \text{ and } v \in s\}$
- Attributes of relations  $r$  and  $s$  should be disjoint
- If attributes are not disjoint, renaming should be used
- Schema is changed
- Applying  $\times$  on

A	B	and	C	D	E
1	1		1	2	7
2	2		2	6	8
			5	7	9

returns

# Cartesian Product in Relational Algebra

- $r \times s = \{\langle u, v \rangle | u \in r \text{ and } v \in s\}$
- Attributes of relations  $r$  and  $s$  should be disjoint
- If attributes are not disjoint, renaming should be used
- Schema is changed
- Applying  $\times$  on

<table><tr><th>A</th><th>B</th></tr><tr><td>1</td><td>1</td></tr><tr><td>2</td><td>2</td></tr></table>		A	B	1	1	2	2	and	<table><tr><th>C</th><th>D</th><th>E</th></tr><tr><td>1</td><td>2</td><td>7</td></tr><tr><td>2</td><td>6</td><td>8</td></tr><tr><td>5</td><td>7</td><td>9</td></tr></table>			C	D	E	1	2	7	2	6	8	5	7	9
A	B																						
1	1																						
2	2																						
C	D	E																					
1	2	7																					
2	6	8																					
5	7	9																					
returns		A	B	C	D	E																	
		1	1	1	2	7																	
		1	1	2	6	8																	
		1	1	5	7	9																	
		2	2	1	2	7																	
		2	2	2	6	8																	
	2	2	5	7	9																		

# From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations

# From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

# From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

```
select title  
from course, offering  
where course.code = offering.coursecode and yr = 2018;
```

# From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

```
select title  
from course, offering  
where course.code = offering.coursecode and yr = 2018;
```

- When two relations contain attributes of the same name, qualification is needed to remove ambiguity
- Example: Find roll number of students in B.Tech. program



- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

```
select title  
from course, offering  
where course.code = offering.coursecode and yr = 2018;
```

- When two relations contain attributes of the same name, qualification is needed to remove ambiguity
- Example: Find roll number of students in B.Tech. program

```
select student.roll  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.';
```

# Select in Relational Algebra

- $\sigma_p(r) = \{t | t \in r \text{ and } p(t)\}$
- $p$  is called the **selection predicate**
- Select all tuples from  $r$  that satisfies the predicate  $p$
- Schema is

# Select in Relational Algebra

- $\sigma_p(r) = \{t | t \in r \text{ and } p(t)\}$
- $p$  is called the **selection predicate**
- Select all tuples from  $r$  that satisfies the predicate  $p$
- Schema is not changed
- Applying  $\sigma_{A=B \wedge D > 5}$  on

A	B	C	D
1	1	2	7
1	2	5	7
2	2	9	3
2	2	8	6

returns

# Select in Relational Algebra

- $\sigma_p(r) = \{t | t \in r \text{ and } p(t)\}$
- $p$  is called the **selection predicate**
- Select all tuples from  $r$  that satisfies the predicate  $p$
- Schema is not changed
- Applying  $\sigma_{A=B \wedge D > 5}$  on

	A	B	C	D
	1	1	2	7
	1	2	5	7
	2	2	9	3
	2	2	8	6
	A	B	C	D
returns	1	1	2	7
	2	2	8	6

# Where

- Specifies conditions that the result tuples must satisfy

# Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

# Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

- May use **and**, **or** and **not** to connect predicates

```
select coursecode  
from offering  
where yr = 2018 and instructor = 10;
```

# Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

- May use **and**, **or** and **not** to connect predicates

```
select coursecode  
from offering  
where yr = 2018 and instructor = 10;
```

- Unused clause is equivalent to **where true**

```
select coursecode  
from offering;
```



# Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

- May use **and**, **or** and **not** to connect predicates

```
select coursecode  
from offering  
where yr = 2018 and instructor = 10;
```

- Unused clause is equivalent to **where true**

```
select coursecode  
from offering;
```

- SQL allows **between** operator (includes both)

```
select coursecode  
from offering  
where yr between 2016 and 2018;
```

# Rename Operation

- SQL allows renaming of relations and attributes to remove ambiguity
- Keyword **as** is used
- Example

```
select student.roll as rollnumber  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.';
```

# Rename Operation

- SQL allows renaming of relations and attributes to remove ambiguity
- Keyword **as** is used
- Example

```
select student.roll as rollnumber  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.' ;
```

- Renaming is necessary when the same relation needs to be used twice
- Example: Find names of students whose cpi is greater than that of "ABC"

# Rename Operation

- SQL allows renaming of relations and attributes to remove ambiguity
- Keyword **as** is used
- Example

```
select student.roll as rollnumber  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.';
```

- Renaming is necessary when the same relation needs to be used twice
- Example: Find names of students whose cpi is greater than that of "ABC"

```
select T.name  
from student as T, student as S  
where T.cpi > S.cpi and S.name = 'ABC';
```

- **as** can be omitted by simply stating **student T**

# String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
  - **\_**: matches any character
  - **%**: matches any substring

# String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
  - **\_**: matches any character
  - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

# String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
  - **\_**: matches any character
  - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```

# String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
  - **\_**: matches any character
  - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```

- Example: Find departments with the name “?E”



# String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
  - **\_**: matches any character
  - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```

- Example: Find departments with the name “?E”

```
select *  
from department  
where name like ‘“_E” ’;
```

# Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**

# Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**
  - For display purposes only and has no actual effect
- Example: Order departments by name

# Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**
  - For display purposes only and has no actual effect
- Example: Order departments by name

```
select *  
from department  
order by name;
```

- Use **desc** to obtain tuples in descending order

```
select *  
from department  
order by name desc;
```

# Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**
  - For display purposes only and has no actual effect
- Example: Order departments by name

```
select *  
from department  
order by name;
```

- Use **desc** to obtain tuples in descending order

```
select *  
from department  
order by name desc;
```

- Default is ascending order (**asc**)

# Set Operations

- Operators **union**, **intersect** and **except** correspond to  $\cup$ ,  $\cap$ ,  $-$
- *Eliminates* duplicates
- For multiset operations, i.e., to retain duplicates, use **all** after the operations
- Example: Find names of all faculty members and students

# Set Operations

- Operators **union**, **intersect** and **except** correspond to  $\cup$ ,  $\cap$ ,  $-$
- *Eliminates* duplicates
- For multiset operations, i.e., to retain duplicates, use **all** after the operations
- Example: Find names of all faculty members and students  

```
(select name from faculty)  
union  
(select name from student);
```
- Example: Find names of faculty members not found in students

# Set Operations

- Operators **union**, **intersect** and **except** correspond to  $\cup$ ,  $\cap$ ,  $-$
- *Eliminates* duplicates
- For multiset operations, i.e., to retain duplicates, use **all** after the operations
- Example: Find names of all faculty members and students

```
(select name from faculty)  
union  
(select name from student);
```

- Example: Find names of faculty members not found in students

```
(select name from faculty)  
except  
(select name from student);
```



# Aggregation in Relational Algebra

- Aggregate functions that can be used are *avg*, *min*, *max*, *sum*, *count*
- Can be applied on groups of tuples as well
- Aggregate operation is of the form  $G_1, \dots, G_k \mathcal{G}_{F_1(A_1), \dots, F_n(A_n)}(E)$  where
  - $G_1, \dots, G_k$  is the list of attributes on which to group (may be empty)
  - Each  $F_i$  is an aggregate function that operates on the attribute  $A_i$
- Applying  $\mathcal{G}_{sum(C)}$  on  $r$

A	B	C	
1	1	5	
1	2	5	returns
2	3	5	
2	4	8	

# Aggregation in Relational Algebra

- Aggregate functions that can be used are *avg*, *min*, *max*, *sum*, *count*
- Can be applied on groups of tuples as well
- Aggregate operation is of the form  $G_1, \dots, G_k \mathcal{G}_{F_1(A_1), \dots, F_n(A_n)}(E)$  where
  - $G_1, \dots, G_k$  is the list of attributes on which to group (may be empty)
  - Each  $F_i$  is an aggregate function that operates on the attribute  $A_i$
- Applying  $\mathcal{G}_{sum(C)}$  on  $r$

A	B	C	returns $\frac{sum(C)}{23}$
1	1	5	
1	2	5	
2	3	5	
2	4	8	

# Grouping and Aggregation in Relational Algebra

- First, the tuples are grouped according to  $G_1, \dots, G_k$
- Then, aggregate functions  $F_1(A_1), \dots, F_n(A_n)$  are applied on each group
- Schema changes to  $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying  $AG_{sum(C)}$  on  $r$

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8
3	4	8

returns

# Grouping and Aggregation in Relational Algebra

- First, the tuples are grouped according to  $G_1, \dots, G_k$
- Then, aggregate functions  $F_1(A_1), \dots, F_n(A_n)$  are applied on each group
- Schema changes to  $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying  $AG_{sum(C)}$  on  $r$

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8
3	4	8

returns

A	sum(C)
1	10
2	13
3	8

# Grouping and Aggregation in Relational Algebra

- First, the tuples are grouped according to  $G_1, \dots, G_k$
- Then, aggregate functions  $F_1(A_1), \dots, F_n(A_n)$  are applied on each group
- Schema changes to  $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying  $AG_{sum(C)}$  on  $r$

A	B	C	returns	A	sum(C)
1	1	5		1	10
1	2	5		2	13
2	3	5		3	8
2	4	8			
3	4	8			

- Applying  $AG_{sum(C)}$  on  $r$

A	B	C	returns
1	1	5	
1	2	5	
1	2	4	

# Grouping and Aggregation in Relational Algebra

- First, the tuples are grouped according to  $G_1, \dots, G_k$
- Then, aggregate functions  $F_1(A_1), \dots, F_n(A_n)$  are applied on each group
- Schema changes to  $(G_1, \dots, G_k, F_1(A_1), \dots, F_n(A_n))$
- Applying  $AG_{sum(C)}$  on  $r$

A	B	C
1	1	5
1	2	5
2	3	5
2	4	8
3	4	8

returns

A	sum(C)
1	10
2	13
3	8

- Applying  $AG_{sum(C)}$  on  $r$

A	B	C
1	1	5
1	2	5
1	2	4

returns

A	B	sum(C)
1	1	5
1	2	9

# Aggregate Functions

- Five operations that work on multisets: **avg**, **min**, **max**, **sum**, **count**
- Example: Find the average cpi of students

# Aggregate Functions

- Five operations that work on multisets: **avg**, **min**, **max**, **sum**, **count**
- Example: Find the average cpi of students

```
select avg(cpi)  
from student;
```

- For set operations, use **distinct**
- Example: Find the total number of students



# Aggregate Functions

- Five operations that work on multisets: **avg**, **min**, **max**, **sum**, **count**
- Example: Find the average cpi of students

```
select avg(cpi)  
from student;
```

- For set operations, use **distinct**
- Example: Find the total number of students

```
select count(distinct roll)  
from student;
```

# Grouping

- To apply aggregate operations on separate groups, use **group by**
- The aggregate operator is applied on *each* group separately
- Example: Find the number of students in each department

# Grouping

- To apply aggregate operations on separate groups, use **group by**
- The aggregate operator is applied on *each* group separately
- Example: Find the number of students in each department

```
select dept, count(distinct roll)  
from student  
group by dept;
```

# Grouping

- To apply aggregate operations on separate groups, use **group by**
- The aggregate operator is applied on *each* group separately
- Example: Find the number of students in each department

```
select dept, count(distinct roll)  
from student  
group by dept;
```

- Attributes in **select** clause outside of aggregate functions *must* appear in **group by** list

# Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

# Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

```
select coursecode, avg(grade)
from registration
group by coursecode
having count(roll) >= 5;
```

- The predicate in **having** is applied *after* forming groups whereas the predicate in **where** is applied *before* doing so

# Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

```
select coursecode, avg(grade)
from registration
group by coursecode
having count(roll) >= 5;
```

- The predicate in **having** is applied *after* forming groups whereas the predicate in **where** is applied *before* doing so
- Example: Find average grade in each course of type 4 where number of students is at least 5

# Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

```
select coursecode, avg(grade)
from registration
group by coursecode
having count(roll) >= 5;
```

- The predicate in **having** is applied *after* forming groups whereas the predicate in **where** is applied *before* doing so
- Example: Find average grade in each course of type 4 where number of students is at least 5

```
select coursecode, avg(grade)
from registration, course
where registration.coursecode = course.code and ctype = 4
group by coursecode
having count(roll) >= 5;
```



# Null

- *null* signifies missing or unknown value
- The predicates *is null* and *is not null* can be used to check for null values
- Example: find courses that do not have a webpage

# Null

- *null* signifies missing or unknown value
- The predicates **is null** and **is not null** can be used to check for null values
- Example: find courses that do not have a webpage

```
select code  
from course  
where webpage is null;
```

- *null* signifies missing or unknown value
- The predicates **is null** and **is not null** can be used to check for null values
- Example: find courses that do not have a webpage

```
select code  
from course  
where webpage is null;
```

- Result of expressions involving null evaluate to null
- Comparison with null returns *unknown*
- Uses three-valued logic as relational algebra
- Aggregate functions ignore null
  - **count(\*)** does *not* ignore nulls

# Nested Subqueries

- A query that occurs in the **where** or **from** clause of another query is called a **subquery**
- Entire query is called **outer query** while the subquery is called **inner query** or **nested query**
- Used in tests for set membership, set cardinality, set comparisons

# Nested Subqueries

- A query that occurs in the **where** or **from** clause of another query is called a **subquery**
- Entire query is called **outer query** while the subquery is called **inner query** or **nested query**
- Used in tests for set membership, set cardinality, set comparisons
- Inner query is evaluated for *each tuple* in the outer query

# Nested Subqueries

- A query that occurs in the **where** or **from** clause of another query is called a **subquery**
- Entire query is called **outer query** while the subquery is called **inner query** or **nested query**
- Used in tests for set membership, set cardinality, set comparisons
- Inner query is evaluated for *each tuple* in the outer query
- When a nested query refers to an attribute in the outer query, it is called a **correlated query**

# Set Membership

- Keyword `in` is used for set membership tests
- Example: Find faculty members who have not offered any course

# Set Membership

- Keyword **in** is used for set membership tests
- Example: Find faculty members who have not offered any course

```
select *  
from faculty  
where fid not in (  
    select instructor  
    from offering );
```



# Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
  - select** < attribute list >
  - from** < relation list >
  - where** < predicate or tuple condition >
  - group by** < group attribute list >
  - having** < group condition >
  - order by** < attribute list >

# Insertion

# Insertion

- `insert into ... values` statement

# Insertion

- `insert into ... values` statement
- Example: Create a new student “ABC” with roll 1897 and department 7

# Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

# Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- May omit schema

```
insert into student  
values (1897, ‘‘ABC’’, 7, 0.0);
```

# Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- May omit schema

```
insert into student  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- If value is not known, specify *null*

```
insert into student  
values (1897, ‘‘ABC’’, 7, null);
```

# Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- May omit schema

```
insert into student  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- If value is not known, specify *null*

```
insert into student  
values (1897, ‘‘ABC’’, 7, null);
```

- To avoid *null*, specify schema

```
insert into student(roll , name, dept)  
values (1897, ‘‘ABC’’, 7);
```



## Insertion (contd.)

- Example: Create a course of code 9 for every department with the same type

## Insertion (contd.)

- Example: Create a course of code 9 for every department with the same type

```
insert into course(code, title , webpage, ctype)
select 9, 'New', null, type
from course
where type in (
    select deptid
    from department );
```

- Query is evaluated fully before any tuple is inserted

## Insertion (contd.)

- Example: Create a course of code 9 for every department with the same type

```
insert into course(code, title , webpage, ctype)
select 9, 'New', null, type
from course
where type in (
    select deptid
    from department );
```

- Query is evaluated fully before any tuple is inserted
- Otherwise, infinite insertion happens for queries like

```
insert into r
select * from r;
```

# Deletion

- delete from ... where statement

# Deletion

- `delete from ... where` statement
- Example: Delete student with roll number 1946

# Deletion

- **delete from ... where** statement
- Example: Delete student with roll number 1946  

```
delete from student  
where roll = 1946;
```
- **where** selects tuples that will be deleted

# Deletion

- `delete from ... where` statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```

- `where` selects tuples that will be deleted
- If `where` is empty,

# Deletion

- `delete from ... where` statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```

- `where` selects tuples that will be deleted
- If `where` is empty, all tuples are deleted



# Deletion

- **delete from ... where** statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```

- **where** selects tuples that will be deleted
- If **where** is empty, all tuples are deleted
- Delete all students

```
delete from student;
```

## Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

```
delete from student  
where cpi < (  
    select avg(cpi)  
    from student );
```

## Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

```
delete from student  
where cpi < (  
    select avg(cpi)  
    from student );
```

- Average is computed before any tuple is deleted
- It is *not* re-computed

## Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

```
delete from student  
where cpi < (  
    select avg(cpi)  
    from student );
```

- Average is computed before any tuple is deleted
- It is *not* re-computed
- Otherwise, average keeps changing
- Ultimately, only the student with the largest CPI remains

# Updating

- `update ... set ... where` statement
- `where` selects tuples that will be updated

# Updating

- `update ... set ... where` statement
- `where` selects tuples that will be updated
- Example: Update value of grade 'E' to 2

# Updating

- **update ... set ... where** statement
- **where** selects tuples that will be updated
- Example: Update value of grade 'E' to 2

```
update grade  
set value = 2.0  
where gradecode = 'E';
```

# Updating

- **update ... set ... where** statement
- **where** selects tuples that will be updated
- Example: Update value of grade 'E' to 2

```
update grade  
set value = 2.0  
where gradecode = 'E';
```

- If **where** is empty, all tuples are updated with the new value
- Example: Increase CPI of all students by 5%



# Updating

- **update ... set ... where** statement
- **where** selects tuples that will be updated
- Example: Update value of grade 'E' to 2

```
update grade  
set value = 2.0  
where gradecode = 'E';
```

- If **where** is empty, all tuples are updated with the new value
- Example: Increase CPI of all students by 5%

```
update student  
set cpi = cpi * 1.05;
```

## Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

## Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important

## Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important
- **case** statement handles conditional updates in a better manner and is sometimes necessary

## Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important
- **case** statement handles conditional updates in a better manner and is sometimes necessary
- Example: Increase CPI of all students by 10% where CPI is less than 6.0, by 5% when less than 8.0, and 2% otherwise

## Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;  
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important
- **case** statement handles conditional updates in a better manner and is sometimes necessary
- Example: Increase CPI of all students by 10% where CPI is less than 6.0, by 5% when less than 8.0, and 2% otherwise

```
update student  
set cpi =  
  case (cpi)  
    when cpi < 6.0 then cpi * 1.10  
    when cpi < 8.0 then cpi * 1.05  
    else cpi * 1.02  
end;
```

# Join in Relational Algebra

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is

# Join in Relational Algebra

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as  $r \times s$  but (potentially) less number of tuples
- The above form is the most general, called the **theta join**



# Join in Relational Algebra

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as  $r \times s$  but (potentially) less number of tuples
- The above form is the most general, called the **theta join**
- **Equality join**: When the join condition only contains equality
  - $r \bowtie_{B=C} s$

# Join in Relational Algebra

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as  $r \times s$  but (potentially) less number of tuples
- The above form is the most general, called the **theta join**
- **Equality join**: When the join condition only contains equality
  - $r \bowtie_{B=C} s$
- **Natural join**: If two relations share an attribute (also its *name*), equality join on that common attribute
  - Denoted by  $*$  or simply  $\bowtie$  without any predicate
  - Changes schema by retaining *only one copy* of common attribute
  - $r * s = r \bowtie s = r \bowtie_{r.A=s.A} s$
- Applying  $\bowtie$  on

A	B		A	C	
1	1	and	1	2	returns
1	2		2	3	
2	1				

# Join in Relational Algebra

- $r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$
- Join is too common a query to not have its own operator
- Schema is same as  $r \times s$  but (potentially) less number of tuples
- The above form is the most general, called the **theta join**
- **Equality join**: When the join condition only contains equality
  - $r \bowtie_{B=C} s$
- **Natural join**: If two relations share an attribute (also its *name*), equality join on that common attribute
  - Denoted by  $*$  or simply  $\bowtie$  without any predicate
  - Changes schema by retaining *only one copy* of common attribute
  - $r * s = r \bowtie s = r \bowtie_{r.A=s.A} s$
- Applying  $\bowtie$  on

A	B		A	C		A	B	C
1	1		1	2	returns	1	1	2
1	2	and	2	3		1	2	2
2	1					2	1	3

# Join

- Join types: **inner join** or simply **join**
- Join conditions: **natural**, **on**  $\langle$  predicate  $\rangle$ , **using** ( $\langle$  attribute list  $\rangle$ )
- Examples

```
student inner join program on student.roll = program.roll;
```

# Join

- Join types: **inner join** or simply **join**
- Join conditions: **natural**, **on**  $\langle$  predicate  $\rangle$ , **using** ( $\langle$  attribute list  $\rangle$ )
- Examples

```
student inner join program on student.roll = program.roll;  
student natural join program;
```

# Join

- Join types: **inner join** or simply **join**
- Join conditions: **natural**, **on**  $\langle$  predicate  $\rangle$ , **using**  $\langle$  attribute list  $\rangle$
- Examples

```
student inner join program on student.roll = program.roll;  
student natural join program;  
student join program using (roll);
```

# Join

- Join types: **inner join** or simply **join**
- Join conditions: **natural**, **on**  $\langle$  predicate  $\rangle$ , **using** ( $\langle$  attribute list  $\rangle$ )
- Examples

```
student inner join program on student.roll = program.roll;  
student natural join program;  
student join program using (roll);
```

- Multiple relations can be joined

```
student join program join registration;
```

# Variants in SQL

- SQL standards have evolved a lot over the years
- Different vendors provide different flavors and may not implement every feature